

Model Selection

Bayesian Optimization for model selection to maximize accuracy.

```
In [1]: from src.classifier.ThreadedBayesianSearcher import  
from src.classifier.BayesianTabularModelSearch import  
from src.classifier.CustomTabularModel import CustomTabularModel  
from src.recommender.BayesianRecommender import BayesianRecommender  
from bayes_opt import BayesianOptimization  
from bayes_opt import UtilityFunction  
import numpy as np  
  
import matplotlib.pyplot as plt  
from matplotlib import gridspec  
%matplotlib inline
```

In [2]:

```

bayesian_optimizer = BayesianSearcher(20)
model = CustomTabularModel(0.5, False, 1000, {'lay
bayesian_optimizer.run_optimization(model, {'laye
bayesian_optimizer.run_optimization(model, {'laye
bayesian_optimizer.run_optimization(model, {'laye

```

Epoch 4: early stopping

 20.00% [4/20]

00:00<00:02]

| epoch | train_loss | valid_loss | accuracy | time |
|-------|------------|------------|----------|-------|
| 0 | 0.689145 | 0.675664 | 0.822695 | 00:00 |
| 1 | 0.681097 | 0.674108 | 0.822695 | 00:00 |
| 2 | 0.679540 | 0.670044 | 0.822695 | 00:00 |
| 3 | 0.679220 | 0.663594 | 0.822695 | 00:00 |

 100.00% [14/14]

00:00<00:00]

Epoch 4: early stopping

 20.00% [4/20]

In [3]:

```
def maximization_function(**params: dict):
    """
    The function whose value we want to maximize.

    :param params: The parameters to set to the model.
    :return: The value of the metric used by the model.
             Expected to be the validation accuracy.
    """
    params['early_stopping'] = False
    model.reset_params(params)
    return model.train(epochs=20, k=1)
```

In [4]:

```
x = [{ 'layer1': i } for i in range(2, 400, 50)]
y = [maximization_function(**layer_size) for layer_size in x]
```

| | | | | |
|---|----------|----------|----------|-------|
| 0 | 0.670903 | 0.663224 | 0.820331 | 00:00 |
| 1 | 0.696800 | 0.660868 | 0.820331 | 00:00 |
| 2 | 0.688626 | 0.651283 | 0.820331 | 00:00 |
| 3 | 0.677082 | 0.635440 | 0.820331 | 00:00 |

 100.00% [14/14]
00:00<00:00]

Epoch 4: early stopping

 20.00% [4/20]
00:00<00:03]

| epoch | train_loss | valid_loss | accuracy | time |
|-------|------------|------------|----------|-------|
| 0 | 0.670903 | 0.663224 | 0.820331 | 00:00 |
| 1 | 0.647447 | 0.659451 | 0.820331 | 00:00 |
| 2 | 0.644057 | 0.650343 | 0.820331 | 00:00 |
| 3 | 0.641955 | 0.634390 | 0.820331 | 00:00 |

In [5]:

```
x_i = list(range(0, len(x)))
```

```
plt.figure(figsize=(10,10))
```

```
plt.title("Single Layer Model Accuracy")
```


```
plt.ylim((0, 1))
```

```
plt.xlabel("Models Tested")
```

```
plt.ylabel("Val Accuracy")
```

```
plt.plot(x_i, y)
```

```
[<matplotlib.lines.Line2D at 0x1a2888c828>]
```



In [6]:

```

def posterior(optimizer, x_obs, y_obs, grid):
    optimizer._gp.fit(x_obs, y_obs)

    mu, sigma = optimizer._gp.predict(grid, return
    return mu, sigma

def plot_gp(optimizer, x, y):
    fig = plt.figure(figsize=(10, 10))
    steps = len(optimizer.space)
    fig.suptitle(
        'Gaussian Process and Utility Function After
        fontdict={'size': 30}
    )

    gs = gridspec.GridSpec(2, 1, height_ratios=[3,
    axis = plt.subplot(gs[0])
    acq = plt.subplot(gs[1])

    x_obs = np.array([res["params"]["layer1"] for res in opt
    y_obs = np.array([res["target"] for res in opt

    mu, sigma = posterior(optimizer, x_obs, y_obs,
    axis.plot(x, y, linewidth=3, label='Target')
    axis.plot(x_obs.flatten(), y_obs, 'D', markersize=
    axis.plot(x, mu, '--', color='k', label='Predicted')

    axis.fill(np.concatenate([x, x[:, :-1]]),
               np.concatenate([mu - 1.9600 * sigma,
               alpha=.6, fc='c', ec='None', label='

```


```
axis.set_xlim((-2, 10))
axis.set_ylim((None, None))
axis.set_ylabel('f( layer1 )', fontdict={'size': 20})
axis.set_xlabel('layer1', fontdict={'size': 20})

utility_function = UtilityFunction(kind="ucb",
utility = utility_function.utility(x, optimize
acq.plot(x, utility, label='Utility Function',
acq.plot(x[np.argmax(utility)], np.max(utility)
        label=u'Next Best Guess', markerfacecolor='red',
acq.set_xlim((-2, 10))
acq.set_ylim((0, np.max(utility) + 0.5))
acq.set_ylabel('Utility', fontdict={'size': 20})
acq.set_xlabel('layer1', fontdict={'size': 20})


axis.legend(loc=2, bbox_to_anchor=(1.01, 1), border=1)
acq.legend(loc=2, bbox_to_anchor=(1.01, 1), border=1)
```

```
In [7]: bayesian_optimizer.optimizer.maximize(init_points=0, r
ot_gp(bayesian_optimizer.optimizer, np.array(x_i).re
```

| | | | | |
|---|----------|----------|----------|-------|
| 4 | 0.698687 | 0.701711 | 0.749409 | 00:00 |
| 5 | 0.691048 | 0.675212 | 0.794326 | 00:00 |
| 6 | 0.675027 | 0.639234 | 0.806147 | 00:00 |
| 7 | 0.660218 | 0.596379 | 0.813239 | 00:00 |
| 8 | 0.654463 | 0.558682 | 0.815603 | 00:00 |
| 9 | 0.637534 | 0.538710 | 0.815603 | 00:00 |

 100.00% [14/14]
00:00<00:00]

Epoch 10: early stopping

 30.00% [6/20]
00:01<00:03]

| epoch | train_loss | valid_loss | accuracy | time |
|-------|------------|------------|----------|-------|
| 0 | 0.772114 | 0.681974 | 0.747045 | 00:00 |
| 1 | 0.741633 | 0.673312 | 0.751633 | 00:00 |

Recommendation

Via Bayesian Optimization, find the params that need to change to reduce the likelihood someone experience depression.

In [25]:

```

from src.recommender.JSONParamReader import JSONPa
import pprint
import warnings
from pandas.io.json import json
import pandas as pd
from bayes_opt import BayesianOptimization, Events
from fastai.basic_data import DataBunch, Tensor
import numpy as np
from src.classifier.CustomTabularModel import Cust
from src.data.DataCsvInterface import DataCsvInter
from src.recommender.JSONParamReader import JSONPa

warnings.simplefilter(action='ignore', category=Fu

bayesian_optimizer = BayesianRecommender()
# Init the model with the best params
model = CustomTabularModel(0.5, False, 1000, {'lay
best_params = JSONParamReader('classifier/logs').g
model.reset_params(best_params)
pp = pprint.PrettyPrinter(indent=4)
print("          Using Model Architecture")
pp.pprint( best_params)

```

Using Model Architecture

```

{'dropout': 0.4170220047, 'layer1': 288.409472883
4}

```


In [26]:

```

data = model.input_data.train_ds[0][0]
data_parsed = []
for element in data.data:
    if type(element) is Tensor:
        data_parsed += list(element.numpy())
    else:
        data_parsed += element

data_init = {key: data_parsed[i] for i, key in enumerate(data.names, model.column_range)}
cr = bayesian_optimizer.get_ranges(data.names, model.column_range)
column_range = {key: cr[key] for key in cr if key in data_init}
model.train(90)

```

4.44% [4/90]

00:01<00:24]

| epoch | train_loss | valid_loss | accuracy | time |
|-------|------------|------------|----------|-------|
| 0 | 0.675767 | 0.658606 | 0.846808 | 00:00 |
| 1 | 0.664987 | 0.645669 | 0.846808 | 00:00 |
| 2 | 0.654010 | 0.630003 | 0.846808 | 00:00 |
| 3 | 0.644451 | 0.604866 | 0.846808 | 00:00 |

100.00% [8/8]

00:00<00:00]

Epoch 4: early stopping

0.8468084931373596

```
In [27]: gaussian_optimizer.run_optimization(model, data_init, (

Value to maximize: 0.4999523162841797
Value to maximize: 0.5003073215484619
Value to maximize: 0.500303328037262
Value to maximize: 0.5007879734039307
Value to maximize: 0.5008745789527893
Value to maximize: 0.5008864402770996
Keeping results: {'target': 0.5008864402770996, 'p
arams': {'bodyweight_categorical': 1.0850132308450
688, 'depressed_categorical': 2.9361659231999284,
'edu_level_categorical': 3.722764174374958, 'emplo
yment_categorical': 1.3035851241440028, 'friends':
3.9945746100373665, 'gender_categorical': 1.0, 'in
come_categorical': 1.0870035421055428, 'income_flo
at': 1.3967942175852635, 'pay_for_sex_categorica
l': 1.0466278587617235, 'prostitution_legal_catego
rical': 1.0768838754090093, 'race_categorical': 1.
0084664602164604, 'sexualallity_categorical': 1.6822
567688459231, 'social_fear_categorical': 3.5620199
16620608, 'virgin_categorical': 1.033908755094694
4}}
[{'target': 0.5008864402770996, 'params': {'bodywe
ight_categorical': 1.0850132308450688, 'depressed_
categorical': 2.9361659231999284, 'edu_level_categ
orical': 3.722764174374958, 'employment_categorica
l': 1.3035851241440028, 'friends': 3.9945746100373
665, 'gender_categorical': 1.0, 'income_categorica
l': 1.0870035421055428, 'income_float': 1.39679421
75852635, 'pay_for_sex_categorical': 1.04662785876
17235, 'prostitution_legal_categorical': 1.0768838
754090093, 'race_categorical': 1.0084664602164604,
'sexualallity_categorical': 1.6822567688459231, 'soc
```

```
ial_fear_categorical': 3.562019916620608, 'virgin_  
categorical': 1.0339087550946944}}]
```

In [64]:

```

after_data = bayesian_optimizer.results[0]['params']

# Sort both dictionaries by keys
data_init = {i:data_init[i] for i in sorted(data_init.keys())}
after_data = {i:after_data[i] for i in sorted(after_data.keys()) if i != 1}

# Exclude changing your sexuality lol
del data_init['sexuality_categorical']
del after_data['sexuality_categorical']


# Before Data
feature_data = [data_init[_] for _ in data_init]
feature_pos = list(range(len(feature_data)))
objects = [_ for _ in data_init]
# Fill in any missing data for running predictions
for key in data_init:
    if key not in after_data:
        print("Missing: " + str(key))
        after_data[key] = data_init[key]

after_feature_data = [after_data[_] for _ in after_data]
after_feature_pos = list(range(len(after_feature_data)))
after_objects = [_ for _ in after_data]

plt.figure(figsize=(10, 10))
b1 = plt.bar(feature_pos, feature_data, align='center')
b2 = plt.bar(after_feature_pos, after_feature_data, align='center')
plt.xticks(feature_pos, objects, rotation=-90)
plt.ylabel('Usage')

```

```
plt.title('Single Individual Recommendation')  
plt.legend((b1[0], b2[0]), ('Before', 'After'))  
plt.show()
```



In [62]:

`data_init`

```
{'age': 0.36457014,
 'bodyweight_categorical': 1,
 'depressed_categorical': 1,
 'edu_level_categorical': 1,
 'employment_categorical': 1,
 'friends': -0.19225857,
 'gender_categorical': 1,
 'income_categorical': 1,
 'income_float': 1.3152382,
 'pay_for_sex_categorical': 1,
 'prostitution_legal_categorical': 1,
 'race_categorical': 1,
 'sexuality_categorical': 1,
 'social_fear_categorical': 1,
 'virgin_categorical': 1}
```

In [78]:

```
changes = np.abs(np.subtract(feature_data, after_f
directions = np.sign(np.subtract(feature_data, aft
```

In [73]:

```
n = 3
max_categories = np.argsort(changes)[-n:][::-1]
```

In []:

`model.input_data.train_dl`

In []:

In []:

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```



```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```